

Understanding and Troubleshooting HPCTOOLKIT Problems

The HPCTOOLKIT Team

03 February 2009

Why isn't my source code found by hpcviewer?

When using `hpcprof` to create performance databases, one must specify the path to the source directories using the `-I dir` flag. If there are multiple source directories, then multiple `-I` flags are required:

```
-I dir1 -I dir2 ... -I dirN
```

The specified list of directory paths, `dir1` through `dirN`, can be either relative or absolute.

In addition, `hpcprof` has a recursive directory search feature that can greatly reduce the number of `-I` flag instances. To conduct a source-file search of a directory *and all of its descendants*, simply append an escaped `'*'` after the last slash, e.g., `/mypath/*` (or `/mypath/'*'`). This greatly simplifies the `hpcprof` command line for a typical project. For example, suppose a project has all of its source in a top-level directory called `src`. Inside `src` are several libraries (like `src/lib1`, `src/lib2`, `src/lib-math`, etc). Then, all of the source is visible to `hpcprof` by using the single flag

```
-I src/'*'
```

NOTE: The `'*'` can be used *only* at the end of a directory path.

For any functions whose source code is not found in the specified directories (e.g., system libraries), `hpcviewer` will generate a synopsis that shows the presence of the function and its line extents (if known), but no source code.

I tried the approach above, but I am still missing source files.

This is a common problem. The cause for this difficulty is that the pathnames associated with performance metrics are extracted from an application binary. The path names encoded in the binary can be

- relative to the directory where the source was compiled, not where you are now as you measure and analyze the code, or
- absolute paths that have encodings for the mount point the directory that existed when the file was compiled, but might be different now (e.g., they were compiled on a different node of a cluster that had a different mount point for the shared file system).

Diagnosing and fixing this problem requires knowing exactly what path names are referenced in the binary and/or perhaps the performance data. Fortunately, this information is supplied by `hpcprof`. If a source file is successfully located, then a

```
msg:    cp:...
```

line appears in the output of `hpcprof`. Unlocated files are deemed 'lost' and there is an output line of the form

```
WARNING:  lost:
```

in the output.

For example, suppose we have an application `app1` whose main source is in a directory `/projs/Apps/app1-src`. The `app1` application is built inside the `app1-src` subdirectory, and it uses source files from a subdirectory `app1-src/special` as well as some source common to all applications, located in `/projs/Apps/common`. When `app1` is built, the `common` source is accessed by relative path `../common`. The `app1` executable is installed on our path.

Now, we switch to our home directory `/h/user/T1` to collect some profile data for `app1`. When we run `hpcstruct` (without the `-I` flag) as follows:

```
hpcstruct -S app1.hpcstruct */*.hpcrun
```

This results in the output

```
msg: Line map : /opt/apps/intel/compilers/10.1/lib/libimf.so
msg: STRUCTURE: /usr/local/bin/app1
msg: Copying source files reached by PATH option to /h/user/T1/hpctoolkit-app1-database
WARNING: lost: app1.c
WARNING: lost: special/xfn1.c
WARNING: lost: ../common/mathx.c
WARNING: lost: ~unknown-file~
WARNING: lost: irc_msg_support.c
```

The `WARNING: lost:` obtains for `~unknown-file~` and `irc_msg_support.c` because these are compiler system files — source is unavailable. The other lost files, however, can be found by using the proper `-I` flag:

```
hpcstruct -I /projs/Apps/'*' -S app1.hpcstruct */*.hpcrun
```

The resulting output:

```
msg: Line map : /opt/apps/intel/compilers/10.1/lib/libimf.so
msg: STRUCTURE: /usr/local/bin/app1
msg: Copying source files reached by PATH option to /h/user/T1/hpctoolkit-app1-database
msg:  cp:/projs/Apps/app1-src/app1.c -> ./projs/Apps/app1-src/app1.c
msg:  cp:/projs/Apps/app1-src/special/xfn1.c -> ./projs/Apps/app1-src/special/xfn1.c
msg:  cp:/projs/Apps/common/mathx.c -> ./projs/Apps/common/mathx.c
WARNING: lost: ~unknown-file~
WARNING: lost: irc_msg_support.c
```

Much better!

Best Practice: Carefully inspect the output of `hpcprof` to determine which files are lost. Then, use the recursive search form of the `-I` flag with absolute paths to the actual source directory.

Why don't the line numbers for loops and/or procedures exactly correspond to what I see in the source code?

To use a cliché, “garbage in, garbage out”. `HPCTOOLKIT` depends on information recorded in the symbol table by the compiler. Line numbers for procedures and loops are inferred by looking at the symbol table information recorded for machine instructions identified as being inside the procedure or loop.

For procedures, often no machine instructions are associated with a procedure's declarations. Thus, the first line in the procedure that has an associated machine instruction is the first line of executable code.

Inlined functions may occasionally lead to confusing data for a procedure. Machine instructions mapped to source lines from the inline function appear in the context of other functions. While `hpcprof`'s methods for handling inline functions are good, some codes can confuse the system.

For loops, the process of identifying what source lines are in a loop is similar to the procedure process: what source lines map to machine instructions inside a loop defined by a backward branch to a loop head. Sometimes compilers sometimes don't properly record the line number mapping.

When the compiler line mapping information is wrong, there is little you can do about it other than to ignore its imperfections, or hand-edit the XML program structure file produced by `hpcstruct`. This technique is used only when truly desperate.